

CS213

Principles of Database Systems(H)

Chapter 12 Index

Shiqi YU 于仕琪

yusq@sustech.edu.cn

Most contents are from Stéphane Faroult's slides

12.1 Index

Shiqi Yu 于仕琪


yusq@sustech.edu.cn

Most contents are from Stéphane Faroult's slides


Problem

SQL

Solution




So let's see what might help find data faster in a table (imagine that you are looking for people by name). It's like with books. If you have few books at home and are looking for one, you can check them one by one. In a bookshop, books need to be ordered by theme. In a big library, they need a very strict organization. Moreover, you cannot rely on any physical order, because some people may look for a book by author, and others by title.



Before computers (and you can still find them), you had drawers where you could look for books by author, title or sometimes subject that were telling you what were the "coordinates" of a book. That's an index.


file 002
block #6
783 bytes from start

} Row locator





We'll see the structure in more detail later, but a database is made of files, themselves usually organized in equal-sized "pages" (or "blocks"). File, block and offset allow to locate very fast any row.

The whole idea of indexing consists in associating values in one or several columns of a table (we may look by groups of columns, such as FIRST_NAME and SURNAME in the PEOPLE table, and want to index them as a combination) to the locator(s) of the row(s) where they can be found.



We build a sorted list of all the values with their locators.





Kafka	locator
Cervantes	locator
Tolstoy	locator
Tolstoy	locator
Twain	locator
Twain	locator
Goethe	locator
Goethe	locator
Alighieri	locator
Shakespeare	locator
Shakespeare	locator
Shakespeare	locator
Dostoyevsky	locator
Dostoyevsky	locator
Dumas	locator

Alighieri

Cervantes

Dostoyevsky

Dostoyevsky

Dumas

Goethe

Goethe

Kafka

Shakespeare

Shakespeare

Shakespeare

Tolstoy

Tolstoy

Twain

Twain

locator

locator


locator

locator


cator

icator

locator



Remember that indexed values are binary values, that will be important!



To search the list easily, we plug a tree above it.

You create an index by giving it a name and specifying table name and column(s)

```
create index <index name>  
on <table name>(<col1>, ... <coln>)
```

Example:

```
create index countries_cont_idx  
on countries(continent)
```

Two columns often queried together can be indexed together; what is indexed is concatenated values (NOT separate values)

```
create index people_surname_born_idx  
on people(surname, born)
```

Composite index

You have actually already created indexes without knowing it:
whenever you declare a PRIMARY KEY or UNIQUE constraint, an
index is created behind your back.

```
create table movies
(
    movieid      integer      not null
        constraint movies_pkey
        primary key,
    title        varchar(100) not null
        constraint "title length"
        check (length(title)::text) <= 100),
    country      char(2)      not null
        constraint movies_country_fkey
        references countries
        constraint "country length"
        check (length(country) <= 2),
    year_released integer      not null
        constraint "year_released numerical"
        check ((year_released + 0) = year_released),
    runtime       integer
        constraint "runtime numerical"
        check ((runtime + 0) = runtime),
    constraint movies_title_country_year_released_key
        unique (title, country, year_released)
);
```


Primary key

INDEX

Unique constraint

It has nothing to do with constraints or the relational theory (with indexes, we are more talking engineering than theory), it's purely practical.

The only way to find quickly in a big table that a supposedly unique value is already recorded is to index the column



Oooops
Already there ...

INSERT

DELETE

Then we should certainly index all columns? Why isn't it done by default? In fact, everything isn't rosy: insertion and deletion always require maintaining table AND indexes. Quite a lot of work.

Table

+ Index



+ Index

+ Index

+ Index

Updating an indexed column isn't only changing its value.


UPDATE



It requires moving things around in the tree, which is more painful work. The location is the same but the key has changed.

Storage

Additionally, indexes use a lot of storage, sometimes more than data! It has a huge impact on operations.



(By "operations" I mean regular activities such as backups)

You can also declare an index to be unique.

```
create unique index <index name>  
on <table name>(<col1>, ... <coln>)
```

Enforces unique
constraint like a
constraint definition

If both are equivalent, then which one
should we use?

Unique index

Unique constraint

Constraints Logic Rules

However (there is always a "however") there are some rare cases when some uniqueness (such as case-insensitive uniqueness in Oracle in a column in which data is in mixed case) cannot be enforced through declarative constraints but can be with the dirty trick of unique indexes.

Constraints, without hesitation. They refer to design.

Indexes Implementation

For naming conventions, some people advocate prefixing an index name with something special.

IDX_MYTABLE_SOMEINDEX

If you have the choice, a suffix is probably a better option (with triggers too). It's possible to list all objects in a database, and, by sorting by name, suffixes allow to see all related objects grouped together.

MYTABLE_SOMEINDEX_IDX

In any case, follow standards that are in place.

FOLLOW NAMING STANDARDS

What matters isn't so much the rule than the fact that everybody is following it and that a name tells you immediately what an object is.


12.2 Performance

Shiqi Yu 于仕琪

yusq@sustech.edu.cn

Most contents are from Stéphane Faroult's slides


0.4
Indexes massively
improve performance.
Scanning a half-million
row table for a unique
value may take a
fraction of a second,
but it will be instant
with an index.



An end-user may not notice much of a difference,
because a sub-second response time is quite acceptable.
It's when the same action is repeated a large number of
times that it makes a difference.

x 100

0.01 40



1 2 3


We are here right in the topic of scalability. Most computer systems are queueing systems. With few users, a mediocre response time is usually OK.



With a lot of queries, later queries will have to wait for earlier queries to be processed before they can be handled. For end-users, the wait time is part of the response time.


Performance

Problems start when the arrival rate of queries becomes faster than the speed at which queries can be processed on average. Then queues lengthen and performance crashes all of a sudden. It may happen on a website that becomes too successful.



The faster
the better.



0.4



Does the super-performance of indexes mean that we should NEVER scan a table? That would be too easy! When we scan a table, the time to scan it is irrespective of how many rows we'll find to return: 1, 100, or several tens of thousands. An index search gives addresses of individual rows.

If we need to retrieve **MANY** rows, there will be a time when plainly scanning the table will be faster.


where column_name in
(select ...)




It arrives earlier than you might think.

What you are using and how you are doing it actually depends on how much data you are retrieving. What is known as "OLTP" (OnLine Transaction Processing) usually makes heavy use of indexes. By contrast, massive batch processes scan a lot, and it's how it should be.


Right algorithm



Right volume



You sometimes find on forums people who are well-wishing but not very experienced who advise people who know hardly less than them to index all of their columns. Don't.



Indexing every column means, once again, that every INSERT will have to write not only into the table but into every index. Ouch.


You always have to compromise. Even if an index makes a search significantly faster, you have to put it in balance with the negative impact on inserts and deletes.






I have carried out tests, checking how many rows I was able to insert in the same amount of time with a variable number of indexes. Here is the result.

Of course you want indexes. You need constraints too. There is nothing wrong with 3 or 4 indexes. But above 5, I begin checking whether they really are beneficial.





**But if we don't want to index ALL columns, the big question becomes:
which ones?**


often used search criteria

There is little need to index a column that you don't use as a search criterion: indexes are primarily here to help you find values faster. In the same way, it doesn't make much sense to index a column especially for the yearly report if it must penalize your inserts all year long.

Typically, in table containing exchange rates, you would index by currency code and date, because one has little meaning without the other. Some people might want to add the exchange rate to the index to find it there without needing to access the table (we may talk more about storage tricks later if we have time)

where column_name = ...

currency_code	as_of_date	exchange_rate
---------------	------------	---------------




Another very important factor is that the column must be **SELECTIVE**, that means that the values it contains are rare and correspond to very few rows. Unique columns and PK are extremely selective by definition. Other columns are a mixed bag.

selective
rare




An example of bad selectivity
would be the word COMPUTER in
the library of the CS department.



It's probably in every book.

Number
of movies



The big problem is with the distribution of values. For some countries (US, Japan, China, India, Britain) we may have many films, and the country isn't selective.

But for other countries we may have one or two films, and then it will be different.

12.3 Execution plan

Shiqi Yu 于仕琪

yusq@sustech.edu.cn

what is the DBMS actually **DOING?**

If there is an index on a column such as COUNTRY, will the DBMS use it? Not necessarily. The optimizer may decide not to use an index. Is there a way to know if it will? Good news: yes.

All DBMS products implement (with slight syntax differences, it won't surprise you) a way to display what is known as the "execution plan", what the optimizer would choose to do to run a query.


explain *<select statement>*



Microsoft®
SQL Server®

set showplan_all on

execution plan




Tables Indexes

In the execution plan, you see tables and indexes that are accessed.

Flickr: Brian Hillegas

explain



Depending on the DBMS, the "explain" command will display the plan directly or populate a table that must be queried.

Many tools can also display graphically an execution plan on demand.

Graphical Environments

SQL Server Management Studio

SQL Developer

IBM Data Studio etc.





Tools Window Community Help

SQLQuery1

```
SELECT          m.title, m.yearreleased
   DISTINCT      credits.[PK_credits_F014496E0E0...]
   FROM movies m
   INNER JOIN credits c
     ON c.movieid = m.movieid
   INNER JOIN people p
     ON p.peopleid = c.peopleid
   WHERE p.surname = 'Bogart'
```

The screenshot shows the Microsoft SQL Server Management Studio (SSMS) interface. A query window titled "SQLQuery1" displays a T-SQL SELECT statement. The execution plan for this query is overlaid on the right side of the window, showing three operations: a "Clustered Index Scan (Clustered)" on the "credits" table, a "Clustered Index Seek (Clustered)" on the "people" table, and a "Clustered Index Seek (Clustered)" on the "movies" table. The cost of the query is listed as 32%.







```
mysql> explain select distinct m.title, m.year_released
-> from movies m
->      inner join credits c
->          on c.movieid = m.movieid
->      inner join people p
->          on p.peopleid = c.peopleid
-> where p.surname = 'Bogart';
```

id select_type table key ref Extra
1 SIMPLE p surname const Using where; Using index
1 SIMPLE c peopleid movies.p.peopleid Using index
1 SIMPLE m PRIMARY movies.c.movieid

3 rows in set (0.00 sec)

```
mysql>
```

Some columns have been
removed



```
movies=# explain select distinct m.title, m.year_released  
movies-# from movies m  
movies-#     inner join credits c  
movies-#         on c.movieid = m.movieid  
movies-#     inner join people p  
movies-#         on p.peopleid = c.peopleid  
movies-# where p.surname = 'Bogart';
```

QUERY PLAN

```
HashAggregate  (cost=5.29..5.30 rows=1 width=222)  
  -> Nested Loop  (cost=2.16..5.28 rows=1 width=222)  
    -> Hash Join  (cost=2.16..4.51 rows=1 width=4)  
      Hash Cond: (c.peopleid = p.peopleid)  
      -> Seq Scan on credits c  (cost=0.00..1.97 rows=97 width=8)  
      -> Hash  (cost=2.15..2.15 rows=1 width=4)  
        -> Seq Scan on people p  (cost=0.00..2.15 rows=1 width=4)  
          Filter: ((surname)::text = 'Bogart'::text)  
  -> Index Scan using movies_pkey on movies m  (cost=0.00..0.76 rows=1 width=226)  
    Index Cond: (movieid = c.movieid)
```

(10 rows)

```
movies=#
```



```
sqlite> explain query plan
...> select distinct m.title, m.year_released
...> from movies m
...>     inner join credits c
...>         on c.movieid = m.movieid
...>     inner join people p
...>         on p.peopleid = c.peopleid
...> where p.surname = 'Bogart';
0|0|2|SEARCH TABLE people AS p USING COVERING INDEX sqlite_autoindex_people_1 (surname=?) (~10)
0|1|1|SEARCH TABLE credits AS c USING AUTOMATIC COVERING INDEX (peopleid=?) (~7 rows)
0|2|0|SEARCH TABLE movies AS m USING INTEGER PRIMARY KEY (rowid=?) (~1 rows)
0|0|0|USE TEMP B-TREE FOR DISTINCT
sqlite>
```

It's worth noting that the same query on the same tables with the same indexes and the same data in the tables may very well result in different execution plans with different DBMS products. Internal algorithms are different, internal data storage is different, some products may be better than others at processing data in a particular way, and of course optimizers are different and may choose different courses.

Beginners often assume that there are "good" execution plans (that only use indexes) and "bad" ones (which scan tables). In fact, it's impossible in most cases to say, by reading two different plans, which one is fastest.




execution plans



The main benefit of execution plans is to check whether the optimizer is more or less doing what you thought it would do.

explain



In particular whether an index is used. The optimizer may choose not to use it. It may also be unable to use it.

Index used?

```
|filmdb=# explain select * from movies where movieid =100;  
                      QUERY PLAN  
-----  
Index Scan using movies_pkey on movies  (cost=0.29..8.30 rows=1 width=31)  
  Index Cond: (movieid = 100)  
(2 rows)
```


```
|filmdb=# explain select * from movies where movieid >100;  
                      QUERY PLAN  
-----  
Seq Scan on movies  (cost=0.00..187.05 rows=9105 width=31)  
  Filter: (movieid > 100)  
(2 rows)
```

12.4 Use Index or Not

Shiqi Yu 于仕琪

yusq@sustech.edu.cn


index key



The impossibility of using an index (or of using it as intended) may come from several reasons. One is with composite indexes, the key of which is concatenated values from several columns. The index can only be used if the lead column(s) appear in the WHERE condition of the query.

The problem is the same as looking up for someone's number in a phone book when you don't know the surname. If you had the surname and address but not the first name, you could do it.

surname	firstname	column 3
---------	-----------	----------



PEOPLE



SURNAME

FIRST_NAME

```
explain select * from people
```

```
where surname = '...'
```

```
and first_name = '...'
```

index can be used


```
where first_name = '...'
```

index can not be used

like '%something%'

Exactly the same problem happens when you are using a LIKE expression that STARTS with a %, or the SUBSTR() function. Without the leading part, no way you can walk the tree that takes you to row addresses.

function()



Actually, the problem is more general than that. Suppose that you have a tree built upon the values that you can find in column C. If you say that $f(C)$ is equal to something, there is no way to find the corresponding value for C in the tree.




Case-sensitive DBMS:

```
where upper(surname) = upper('some input')
```

A typical, and common, problem is to perform, with a case-sensitive DBMS a case insensitive search in a column when data is in mixed case.

A tree is based on the ordering of values. It's how values compare to the values in the node that you are visiting which tells you which subtree you should visit next.



The problem is that the internal codes of letters consider that "a" is greater than "Z" (and don't talk about accented letters)

where upper(surname) = 'MARVIN'

Where to search?

"smaller" values

MILES

O'brien

Stewart

marvin

wayne

"bigger" values

If you imagine that you have a binary tree, find 'Stewart' at the root, and look for something equal to MARVIN when set to uppercase without knowing in which case that something is written, you are toast.

It could actually be anywhere in the tree.

where upper(surname) = 'MARVIN'

Same story with functions that extract date parts.
Use them in a WHERE clause, your index is dead.
You should always express conditions on dates as
range conditions.

~~where extract(month from date_column) = 6
and extract(year from date_column) =
extract(year from current_date)~~




Same story again with implicit conversions.

If they are performed
in the wrong direction,
wave farewell to the index.



and varchar_code = 12345678

```
select cast('12345678' as varbinary(max)) as string_12345678,  
        cast(12345678 as varbinary(max)) as number_12345678;
```



string_12345678	number_12345678
0x3132333435363738	0x00BC614E

Ooops ...

Numeric	String	Date
4	'14'	'25-JUL-1603'
14	'25'	'4-JUL-1776'
25	'4'	'14-JUL-1789'


The reason is simple. Index search is based on ordering (inside the tree) and ordering is different with different datatypes. Convert datatypes, you break the ordering.

```
insert into table_name(column_name)  
values(upper(<input>))
```

Some of these issues can be taken care of by only storing appropriate data in your tables, for instance not storing mixed case.

```
select *  
from people  
where soundex(surname) = soundex('Stuart')
```

In some other cases, though, the search MAY require applying a function to the indexed column.



Is there a way to apply a function to a column and yet benefit of the quick access provided by an index? Some products actually allow indexing an expression (more on that coming soon). Otherwise the answer is yes, by cheating.





people

peopleid	first_name	surname	born	died	surname_soundex
1	James	Stewart	1908	1997	S363
2	Humphrey	Bogart	1899	1957	B263

What you can do for instance is add a column that stores the soundex. This happily violates the rules of good normalization (... it depends on another column that is a part of a key), but then you can index it and apply the search to this new redundant column.

people

Trigger – before Insert / for each row

peopleid	first_name	surname	born	died	surname_soundex
1	James	Stewart	1908	1997	S363
2	Humphrey	Bogart	1899	1957	B263

If you have full control of the program that inserts data into the table you can modify it so as to insert soundex(SURNAME) at the same time as you insert SURNAME. If not, there is (at the cost of far slower uploads) the option of a before insert/for each row trigger that does it on the fly.


S363



```
select *  
from people  
where surname_soundex = soundex('Stuart')
```

Having the column and having indexed it, you no longer need to apply any function to the searched column, and everything is fine.

Most products actually allow you to do something cleaner by indexing an expression (sometimes called "generated" or "virtual" column), which can be the result of a function.



$f()$


It only works if the expression or function is

Deterministic

transformation

A large black italicized $f()$ is positioned on the left. To its right, a red arrow points upwards from the word "transformation" towards the text "It only works if the expression or function is". Below this, the word "Deterministic" is written in a large blue sans-serif font. A red arrow points from the bottom of the $f()$ symbol towards the word "transformation".

Same input



Same output

"Deterministic" means that the same input will always generate the same output. Many commonly used functions aren't deterministic because their result is affected by database settings (localization settings most often).

```
datename(month, '1970-01-01')  
'Januar' 'Enero' 'Janeiro'  
'January' '1月' 'Janvier'
```

This SQL Server function isn't deterministic, because if the DBA changes the language by default of the database, what will be returned will be different. You can imagine the scenario: language is English, you create an index, plenty of "January" stored in the index, the DBA switches the language to Spanish and you search the index for "Enero" ... SQL Server prevents you from indexing this function.



Thursday?

PORTUGAL

Day #5

SPAIN

Day #4

MOROCCO

Day #6

You might think that functions that return numbers are safer. Conventions for day numbering actually vary with countries.

Map based on a map by Joan M. Boràs

`upper(column_name)`

Fortunately, there are still many functions that truly are deterministic! If your DBMS allows it, you can index `UPPER(...)` without any problem. Same story with `SOUNDEX(...)`. When your statement will be analyzed the SQL engine will be able to notice your using the function and will use the index.

